

4. Continuous Space Search

Course: Introduction to AI

Instructor: Saumya Jetley

Teaching Assistant(s): Raghav Awasty & Subhrajit Roy

October 14, 2022

What is a continuous space?*



- Space with *uncountably* infinite states
- Connected path from one state to the next

What is a continuous space?*



(Background on problem spaces)

- Problem statements are associated with a *spaces of states* e.g geographical space or parameter space
- Discrete state spaces have discrete number of states (can be infinite)
- Continuous state spaces have an uncountably infinite number of states

What is a continuous space?*




- Is the 3D space we live in discrete or continuous?
- So, a maze (in a park in the real world) is a continuous space?
- NOTE: We can approximate a continuous space with a discrete space by binning parts of it.



- We have a continuous manifold of states
- Keeping the terminology from before, we have:
 - Non-goal states
 - A goal state
- **How do we differentiate the goal state from non-goal states?**



- How do we differentiate the goal state from non-goal states?
- Let's say we put a score on states, given by a score function $f(x)$
- A landscape has both “location” (x) (defined by the state) and “elevation” ($f(x)$), and looks like below:





- How do we differentiate the goal state from non-goal states?
- If elevation/score corresponds to cost, then the goal would be to find the lowest point — a **global minimum**.
- If elevation/score corresponds to performance function, then the goal would be to find the highest peak—a **global maximum**.

Convex vs non-convex functions



- A *set of points* S is *convex* if the line joining any two points in S is also contained within S
- A *convex function* is one for which the space “above” it forms a *convex set*
- By definition, convex functions have no local (as opposed to global) minima




Figure: Example of a convex set (left) and a non-convex set (right)



To reach the goal state - let's say **Maximise an Objective Function**
 $f(x)$

The path to the goal is not important anymore - reaching the goal state is - in reasonable time

State function, not a path function

Different Approaches:

1. Solve as a discrete space search problem
 - Binning
2. Using derivative of $f(x)$ - Solve for $f'(x) = 0$
 - Gradient descent
 - Newton's Method



1. Discrete space search

- Perform binning
 - For each continuous variable in the state space, divide its domain into fixed number of equally spaced bins.
- Use standardised discrete search methods



We could make use of search algorithms such as,

- Hill Climbing
- Simulated Annealing
- Local beam search
- Genetic algorithms



1.1. Hill climbing?

- The search algorithm (**steepest-ascent** version) is a loop that continually moves in the direction of steepest increasing value i.e. *fastest uphill*.
- Terminates when reaches a “peak” - no neighbor has higher value.


Algorithm 1 Algorithm for Hill climbing (Steepest Ascent)

```
1:  $i \leftarrow \text{initialstate}$ 
2: repeat
3:    $\text{BestScore} = f(i)$ ,  $\text{BestState} = i$ 
4:   while  $N(i)$  is non empty do
5:     generate  $s \in N(i)$ , remove  $s$  from  $N(i)$ 
6:     if  $f(s) > \text{BestScore}$  then
7:        $\text{BestScore} := f(s)$ 
8:        $\text{BestState} := s$ 
9:     end if
10:  end while
11:   $i := \text{BestState}$ 
```

State-space diagram of Hill Climbing



- **Local maxima:** Higher than its neighbours but not highest
- **Global maxima:** Here, the value of the objective function is highest
- **Current State:** Current State is the state where the agent is present



Key challenges for hill climbing



- **Greedy Approach:** The search moves in the direction of optimizing the cost i.e. finding local maxima/minima
- **No Backtracking:** It cannot remember the previous state of the system so backtracking to the previous state is not possible
 - **Local Maxima:** The algorithm terminates when the current node is a local maximum but there are better solutions.
 - **Ridge:** Imagine walking along the ridge of a mountain to reach the top - many local maxima along the way!
 - **Plateau:** Plateau is the region where all the neighbouring nodes have the same value - possibility of getting lost!



- Very useful in routing-related problems - Travelling Salesmen Problem, Job Scheduling, Chip Designing, Portfolio Management
- Good in solving the optimization problem using limited computation power
- 14% success rate with vanilla version
- 94% success rate when lateral moves are permitted*
- Extension to hill climbing
 - Stochastic hill climbing
 - First-choice hill climbing
 - Random restart hill climbing



1.1.a. Stochastic Hill Climbing

- **Stochastic hill climbing** chooses at random from among the uphill moves; the probability of selection can vary with the steepness of the uphill move.
- Usually converges more slowly, but in some state landscapes, it finds better solutions.

Algorithm 2 Algorithm for Stochastic Hill climbing

```
1:  $i \leftarrow \text{initialstate}$ 
2: while  $t \leq \text{maxIter}$  do
3:   Generate all  $s \in N(i)$ 
4:   Select  $s$  with probability
```

$$p(s) = \frac{1}{1 + e^{(\text{score}(i) - \text{score}(s))}}$$

```
5:    $i \leftarrow s$  &  $t \leftarrow t + 1$ 
```

Reference: A. Mazidi M. Fakhrahmad M. Sadreddini. (2016). A meta-heuristic approach to CVRP problem: local search optimization based on GA and ant colony. JACR.

1.1.b. First-Choice Hill Climbing



- **First-choice hill climbing** implements stochastic hill climbing by generating successors randomly until one is generated that is better than the current state.
- Choose the first randomly generated *uphill* move!
- This is a good strategy when a state has many (e.g., thousands) of successors.
- Greedy, incomplete, and sub-optimal

1.1.c. Random-restart Hill Climbing*



- It conducts a series of hill-climbing searches from randomly generated initial states, until a goal is found.
- It is trivially complete with probability approaching 1, because it will eventually generate a goal state as the initial state.
- Stops each search at local maxima - or better yet, stops after a fixed amount of time.

1.2. Simulated annealing



- A hill-climbing algorithm that never makes “downhill” moves toward states with lower value (or higher cost) is guaranteed to be incomplete, because it can get stuck on a local maximum
- Escapes local maxima by allowing bad moves (by gradually decreasing their size and frequency)
- **Simulated Annealing** combine hill climbing with a random walk in some way that yields both efficiency and completeness

1.2. Simulated annealing - Algorithm



Algorithm 3 Algorithm for Simulated annealing

```
1:  $i \leftarrow \text{initialstate}$ 
2: while  $t \leq \text{maxIter}$  do
3:   Generate  $s \in N(i)$ 
4:   if  $\text{score}(i) < \text{score}(s)$  then
5:      $i := s$ 
6:   else
7:     Select  $s$  with probability
```

$$p(s) = \frac{1}{1 + e^{\frac{(\text{score}(i) - \text{score}(s))}{T}}}$$

```
8:    $i \leftarrow s$  &  $t \leftarrow t + 1$ 
```



1.3. Local Beam Search

- The **local beam search** algorithm keeps track of k states rather than just one.
- *Here, useful information is passed among the parallel search threads.* (main difference with random-restart search)

Algorithm 4 Algorithm for Local Beam Search

function: BEAM-SEARCH(problem, k) **return** a solution state with k randomly generated states

repeat

 generate all successors of k states

if any of them is a solution **then** return it

else

 select the k best successors

1.3. Local Beam Search - Algorithm



Characteristics:

- Draws all threads to a common region of goodness
- Can suffer from lack of diversity (get concentrated in small region of state-space)
- A variant called **stochastic beam search**, analogous to stochastic hill climbing, chooses k successors at random, with the probability of choosing a given successor being an increasing function of its value.

1.4. Genetic Algorithms



- A **genetic algorithm** (or GA) is a variant of stochastic beam search in which successor states are generated by combining two parent states rather than by modifying a single state.
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Evaluation function (fitness function) \mapsto Higher = better

1.4. Genetic Algorithms



■ Pros:

- Random exploration can find solutions that local search cannot (via crossover primarily)
- Appealing connection to human evolution

■ Cons:

- Large number of “tunable” parameters
- Difficult to replicate performance from one problem to another
- Lack of good empirical studies comparing to simpler methods

1.4. Genetic Algorithm



Algorithm 5 Genetic Algorithm

```
function:  GENETIC-ALGORITHM(population, fitness-function)
return an individual
repeat
    initialize with new population with  $\emptyset$ 
    for  $i = 1$  to  $\text{size}(\text{population})$  do
         $x = \text{random-select}(\text{population}, \text{fitness-function})$ 
         $\text{child} = \text{cross-over}(x, y)$ 
        mutate ( $\text{child}$ ) with small random probability
        add  $\text{child}$  to new population
    end for
     $\text{population} = \text{new-population}$ 
until some individual is fit enough or enough time has elapsed return
the best individual in population w.r.t. the fitness function  $= 0$ 
```



- 1 How to choose the factor for binning?
- 2 Combinatorially complex when state-space is multidimensional?
- 3 What happens when branching factor b is in thousands?
- 4 What happens when branching factor is infinite?

Does it serve us better to work in the continuous space or find a continuous space approximation?



2. Using the derivative

- Remember, when $f'(x) = 0$, the $f(x)$ is
 - at a maximum
 - at a minimum
 - at a point of inflection
- Solve $f'(x) = 0$
- So really we have to find the roots of $f'(x)$
- Then, check each x in $f(x)$ for a maximum (or check $f''(x)$ for each x)
 - $f''(x) = 0$, then point of inflection
 - $f''(x) = \text{positive}$, then minimum
 - $f''(x) = \text{negative}$, then maximum

2. Using the derivative




What if you can't solve $f'(x) = 0$ but you can still figure out the derivative at each point?

- I. Use Newton's Method
- II. Use Gradient descent, or



2.i. Newton's Method of Root Finding

- We are trying to find the root of some function $g(x)$
 - Remember, this is the derivative of the function that we want to optimise ($g(x) = f'(x)$)
- Assume $g(x)$ is linear, and
- You know the value of $g(x)$ and $g'(x)$ for your initial guess x_0



2.i. Newton's Method of Root Finding




Algorithm 6 Algorithm for Newton's root finding

Require: $g(x), g'(x)$

- 1: $x_0 \leftarrow \text{initialguess}$
 - 2: **while** $(x_{i+1} - x_i) \geq \epsilon$ **do**
 - 3: $x_{i+1} = x_i - g(x_i)/g'(x_i)$
 - 4: **end while**
-



2.ii. Gradient Descent (Ascent)





2.ii. Gradient Descent (Ascent)

- Pick an initial value for x_0
- Pick a step size (also called - learning rate lr)

$$x_{i+1} = x_i - lr \frac{d(f(x))}{dx}$$

- Step size too small (takes forever to reach maximum)
- Step size too large (“step over” the maximum)
- Repeat until the sign of $f'(x)$ changes or $f'(x) = 0$



How to fix the step size?

- Variable step size scheme - I
 - Start with large step size and find range within which maximum lies
 - Repeat using smaller step size inside that range
 - Continue until you have enough precision
- Variable step size scheme - II
 - Start by using a small step size
 - As long as the sign of $f'(x)$ doesn't change, double the step size
 - When sign changes, drop back to initial step size and continue search



- **Chapter 4: Artificial Intelligence, A Modern Approach (3rd Edition)** - Stuart Russell and Peter Norvig



Overview

1 Introduction

- What is a continuous space?
- Discrete vs. continuous space

2 Searching in continuous space

- Search landscape: Cost function
- Goal of search

3 Optimisation in Continuous Spaces

- Convex vs. non-convex functions

4 Persisting with discrete search

- Hill Climbing
- Simulated annealing
- Local beam search
- Genetic algorithms

5 Derivatives for search

- Newton's Method of root finding
- Gradient Descent